

**FINAL CHALLENGE****Name:** Manoj Gedela**Kaggle Username:** mg56648

Leaderboard/Display name: Achilles last stand

**Values:**

<b>Model</b>	<b>Kaggle Test AUC(private score)</b>	<b>Comments</b>
Neural network	0.8295167	Kernel_initializer = 'normal'; activation = 'relu'; Activation("sigmoid"), Dense(48), epochs = 3, batch_size = 64
Neural network with embedding	0.937338	Same parameters as above but with embedding
<b>xgboost</b>	<b>0.95216</b>	<ul style="list-style-type: none"><li>• Train extra 2 was created with train_extra, train_sample and some selected rows in the dataset</li><li>• parameter tuning was done to find out the optimal parameters</li><li>• These optimal parameters was used in the final code</li></ul>

**Best Model:** xgboost**Best AUC:** 0.95216

## Submissions: Username : Achilles Last stand

18 submissions for <a href="#">Achilles last stand</a>		Sort by <a href="#">Private Score</a>	
All	Successful	Selected	
Submission and Description	Private Score	Public Score	Use for Final Score
<a href="#">xgboost3 (version 2/2)</a> a day ago by <a href="#">Achilles last stand</a> From "xgboost3" Script	0.9521600	0.9520327	<input type="checkbox"/>
<a href="#">xgboost2 (version 3/3)</a> a day ago by <a href="#">Achilles last stand</a> From "xgboost2" Script	0.9515578	0.9509042	<input type="checkbox"/>
<a href="#">xgboost1 (version 1/1)</a> a day ago by <a href="#">Achilles last stand</a> From "xgboost1" Script	0.9500374	0.9487709	<input type="checkbox"/>
<a href="#">xgb1_withtrain2.csv</a> 2 days ago by <a href="#">Achilles last stand</a> xgb1_withtrain2	0.9500374	0.9487709	<input type="checkbox"/>
<a href="#">xgb1_withtrain1.csv</a> 2 days ago by <a href="#">Achilles last stand</a> xgb1_withtrain1	0.9498081	0.9482843	<input type="checkbox"/>
<a href="#">xgb2.csv</a> 3 days ago by <a href="#">Achilles last stand</a> xgb2	0.9490338	0.9476135	<input type="checkbox"/>
<a href="#">embedding2.csv</a> 4 days ago by <a href="#">Achilles last stand</a> deep neural network 2	0.9373380	0.9390608	<input type="checkbox"/>
<a href="#">sub_nn (1).csv</a> 4 days ago by <a href="#">Achilles last stand</a> embedding	0.8697475	0.8772484	<input type="checkbox"/>
<a href="#">NNprediction_2.csv</a> 4 days ago by <a href="#">Achilles last stand</a> nn2	0.8295167	0.8191874	<input type="checkbox"/>
<a href="#">NNprediction_3.csv</a> 4 days ago by <a href="#">Achilles last stand</a> nn3	0.8048027	0.8004654	<input type="checkbox"/>
<a href="#">NNwithoutembedding.csv</a> 3 days ago by <a href="#">Achilles last stand</a> NN without embedding	0.5030131	0.5030111	<input type="checkbox"/>
<a href="#">NNprediction.csv</a> 4 days ago by <a href="#">Achilles last stand</a> nn1	0.5000000	0.5000000	<input type="checkbox"/>

## CODE SCREENSHOTS:

### Main Code:

```
import numpy as np
import pandas as pd
import math
import time
import xgboost as xgb
from sklearn.cross_validation import train_test_split

#####-----MAIN CODE-----#####

df = pd.read_csv('../input/talkingdata-adtracking-fraud-detection/train.csv', skiprows = 60000000, nrows = 5000000)
df.columns = ['ip', 'app', 'device', 'os', 'channel', 'click_time', 'attributed_time', 'is_attributed']
df1 = df[df['is_attributed'] == 1]
df2 = df[df['is_attributed'] == 0]
del df
df0 = pd.read_csv('../input/talkingdata-adtracking-fraud-detection/train.csv', skiprows = 160000000, nrows = 5000000)
df0.columns = ['ip', 'app', 'device', 'os', 'channel', 'click_time', 'attributed_time', 'is_attributed']
df3 = df0[df0['is_attributed'] == 1]
df4 = df0[df0['is_attributed'] == 0]
del df0
df1 = df1.ix[np.random.random_integers(0, len(df1), 120000)]
df2 = df2.ix[np.random.random_integers(0, len(df2), 240000)]
df3 = df3.ix[np.random.random_integers(0, len(df3), 60000)]
df4 = df4.ix[np.random.random_integers(0, len(df4), 120000)]
df5 = pd.read_csv('../input/talkingdata-adtracking-fraud-detection/train_sample.csv')
df6 = pd.read_csv('../input/train-extra-2/train_extra2.csv')
df = pd.concat([df1, df2, df3, df4, df5, df6])
df = df.drop_duplicates()

def PreProcessTime(df):
    df['click_time'] = pd.to_datetime(df['click_time']).dt.date
    df['click_time'] = df['click_time'].apply(lambda x: x.strftime('%Y%m%d')).astype(int)
    return df

y = train['is_attributed']
train.drop(['is_attributed', 'attributed_time'], axis=1, inplace=True)
store = pd.DataFrame()
store['click_id'] = test['click_id']
test.drop('click_id', axis=1, inplace=True)

params = {'eta': 0.1, 'max_depth': 5, 'subsample': 0.9, 'colsample_bytree': 0.7, 'colsample_bylevel': 0.7, 'min_child_weight': 50,
          'alpha': 4, 'objective': 'binary:logistic', 'eval_metric': 'auc', 'random_state': 238, 'scale_pos_weight': 150,
          'silent': True}

x1, x2, y1, y2 = train_test_split(train, y, test_size=0.1, random_state=238)

totallist = [(xgb.DMatrix(x1, y1), 'train'), (xgb.DMatrix(x2, y2), 'valid')]
model = xgb.train(params, xgb.DMatrix(x1, y1), 400, totallist, maximize=True, verbose_eval=10)

store['is_attributed'] = model.predict(xgb.DMatrix(test), ntree_limit=model.best_ntree_limit)

store.to_csv('xgb1_withtrain5.csv', index=False)
```

Code for tuning to find out the optimal parameters to be used in the final code:

```
from xgboost.sklearn import XGBClassifier
from sklearn import cross_validation, metrics
from sklearn.grid_search import GridSearchCV

#####-----TUNING-----#####

train = pd.read_csv('train_extra.csv')
target = 'is_attributed'

def modelfit(alg, dtrain, predictors, useTrainCV=True, cv_folds=5, early_stopping_rounds=50):

    if useTrainCV:
        xgb_param = alg.get_xgb_params()
        xgtrain = xgb.DMatrix(dtrain[predictors].values, label=dtrain[target].values)
        cvresult = xgb.cv(xgb_param, xgtrain, num_boost_round=alg.get_params()['n_estimators'], nfold=cv_folds,
                           metrics='auc', early_stopping_rounds=early_stopping_rounds)
        alg.set_params(n_estimators=cvresult.shape[0])

    #Fit the algorithm on the data
    alg.fit(dtrain[predictors], dtrain['is_attributed'], eval_metric='auc')

def PreProcessTime(df):
    df['click_time'] = pd.to_datetime(df['click_time']).dt.date
    df['click_time'] = df['click_time'].apply(lambda x: x.strftime('%Y%m%d')).astype(int)
    return df

train = PreProcessTime(train)
train = train.drop(['attributed_time', 'ip'], axis=1)

#Choose all predictors except target
predictors = [x for x in train.columns if x not in [target]]

#Initial values before parameter tuning(Default parameters)
xgb1 = XGBClassifier(learning_rate =0.1, n_estimators=1000, max_depth=5, min_child_weight=1, gamma=0, subsample=0.8,
                     colsample_bytree=0.8, objective= 'binary:logistic', nthread=4, scale_pos_weight=1, seed=27)

modelfit(xgb1, train, predictors)

#---Tune max_depth and min_child_weight---
param_test1 = {'max_depth':list(range(3,10,2)), 'min_child_weight':list(range(1,6,2))}

gsearch1 = GridSearchCV(estimator = XGBClassifier( learning_rate =0.1, n_estimators=1000, max_depth=5, min_child_weight=1,
                                                    gamma=0, subsample=0.8, colsample_bytree=0.8, objective= 'binary:logistic',
                                                    nthread=4, scale_pos_weight=1, seed=27),
                        param_grid = param_test1, scoring='roc_auc', n_jobs=4, iid=False, cv=5)

gsearch1.fit(train[predictors], train[target])

gsearch1.grid_scores_, gsearch1.best_params_, gsearch1.best_score_

#---Tune Gamma---
param_test3 = {'gamma':[i/10.0 for i in range(0,5)]}

gsearch3 = GridSearchCV(estimator = XGBClassifier(learning_rate =0.1, n_estimators=1000, max_depth=7, min_child_weight=1, gamma=0,
                                                    subsample=0.8, colsample_bytree=0.8, objective= 'binary:logistic', nthread=4,
                                                    scale_pos_weight=1, seed=27),
                        param_grid = param_test3, scoring='roc_auc', n_jobs=4, iid=False, cv=5)

gsearch3.fit(train[predictors], train[target])

gsearch3.grid_scores_, gsearch3.best_params_, gsearch3.best_score_
```

```

#---Tune Subsample and colsample_bytree---
param_test4 = {'subsample':[i/10.0 for i in range(6,10)],'colsample_bytree':[i/10.0 for i in range(6,10)]}

gsearch4 = GridSearchCV(estimator = XGBClassifier( learning_rate =0.1, n_estimators=1000, max_depth=4,min_child_weight=2,
                                                    gamma=0.2, subsample=0.8, colsample_bytree=0.8,objective= 'binary:logistic',
                                                    nthread=4, scale_pos_weight=1,seed=27)),
                        param_grid = param_test4, scoring='roc_auc',n_jobs=4,iid=False, cv=5)

gsearch4.fit(train[predictors],train[target])

gsearch4.grid_scores_, gsearch4.best_params_, gsearch4.best_score_

#---Tuning Regularization parameters---

param_test6 = {'reg_alpha':[1e-5, 1e-2, 0.1, 1, 100]}

gsearch6 = GridSearchCV(estimator = XGBClassifier( learning_rate =0.1, n_estimators=1000, max_depth=4, min_child_weight=2,
                                                    gamma=0.2, subsample=0.65, colsample_bytree=0.85, objective= 'binary:logistic',
                                                    nthread=4, scale_pos_weight=1,seed=27)),
                        param_grid = param_test6, scoring='roc_auc',n_jobs=4,iid=False, cv=5)
gsearch6.fit(train[predictors],train[target])
gsearch6.grid_scores_, gsearch6.best_params_, gsearch6.best_score_

#---Tuning scale_pos_weight---
param_test7 = {'scale_pos_weight':[100,200,300,400]}

gsearch7 = GridSearchCV(estimator = XGBClassifier( learning_rate =0.1, n_estimators=1000, max_depth=4, min_child_weight=2,
                                                    gamma=0.2, subsample=0.65, colsample_bytree=0.85, objective= 'binary:logistic',
                                                    nthread=4, scale_pos_weight=1,seed=27)),
                        param_grid = param_test6, scoring='roc_auc',n_jobs=4,iid=False, cv=5)
gsearch7.fit(train[predictors],train[target])
gsearch7.grid_scores_, gsearch6.best_params_, gsearch6.best_score_

```